

Scenario Description

- WVCorp: the company you (the data scientist) work for
 - WVCorp has user forums and discussion boards for each of their products, where customers can discuss issues and features.
 - “Buzz”: when a topic on the user forum has a very high activity level -- considered an indication of user interest in that topic.
- eRead: WVCorp’s ebook reader product
- TimeWrangler: WVCorp’s time-management app
- BookBits: A competitor’s ebook reader product
- GCal: a third-party cloud-based calendar infrastructure that TimeWrangler can integrate with

Using The Buzz Prediction Model

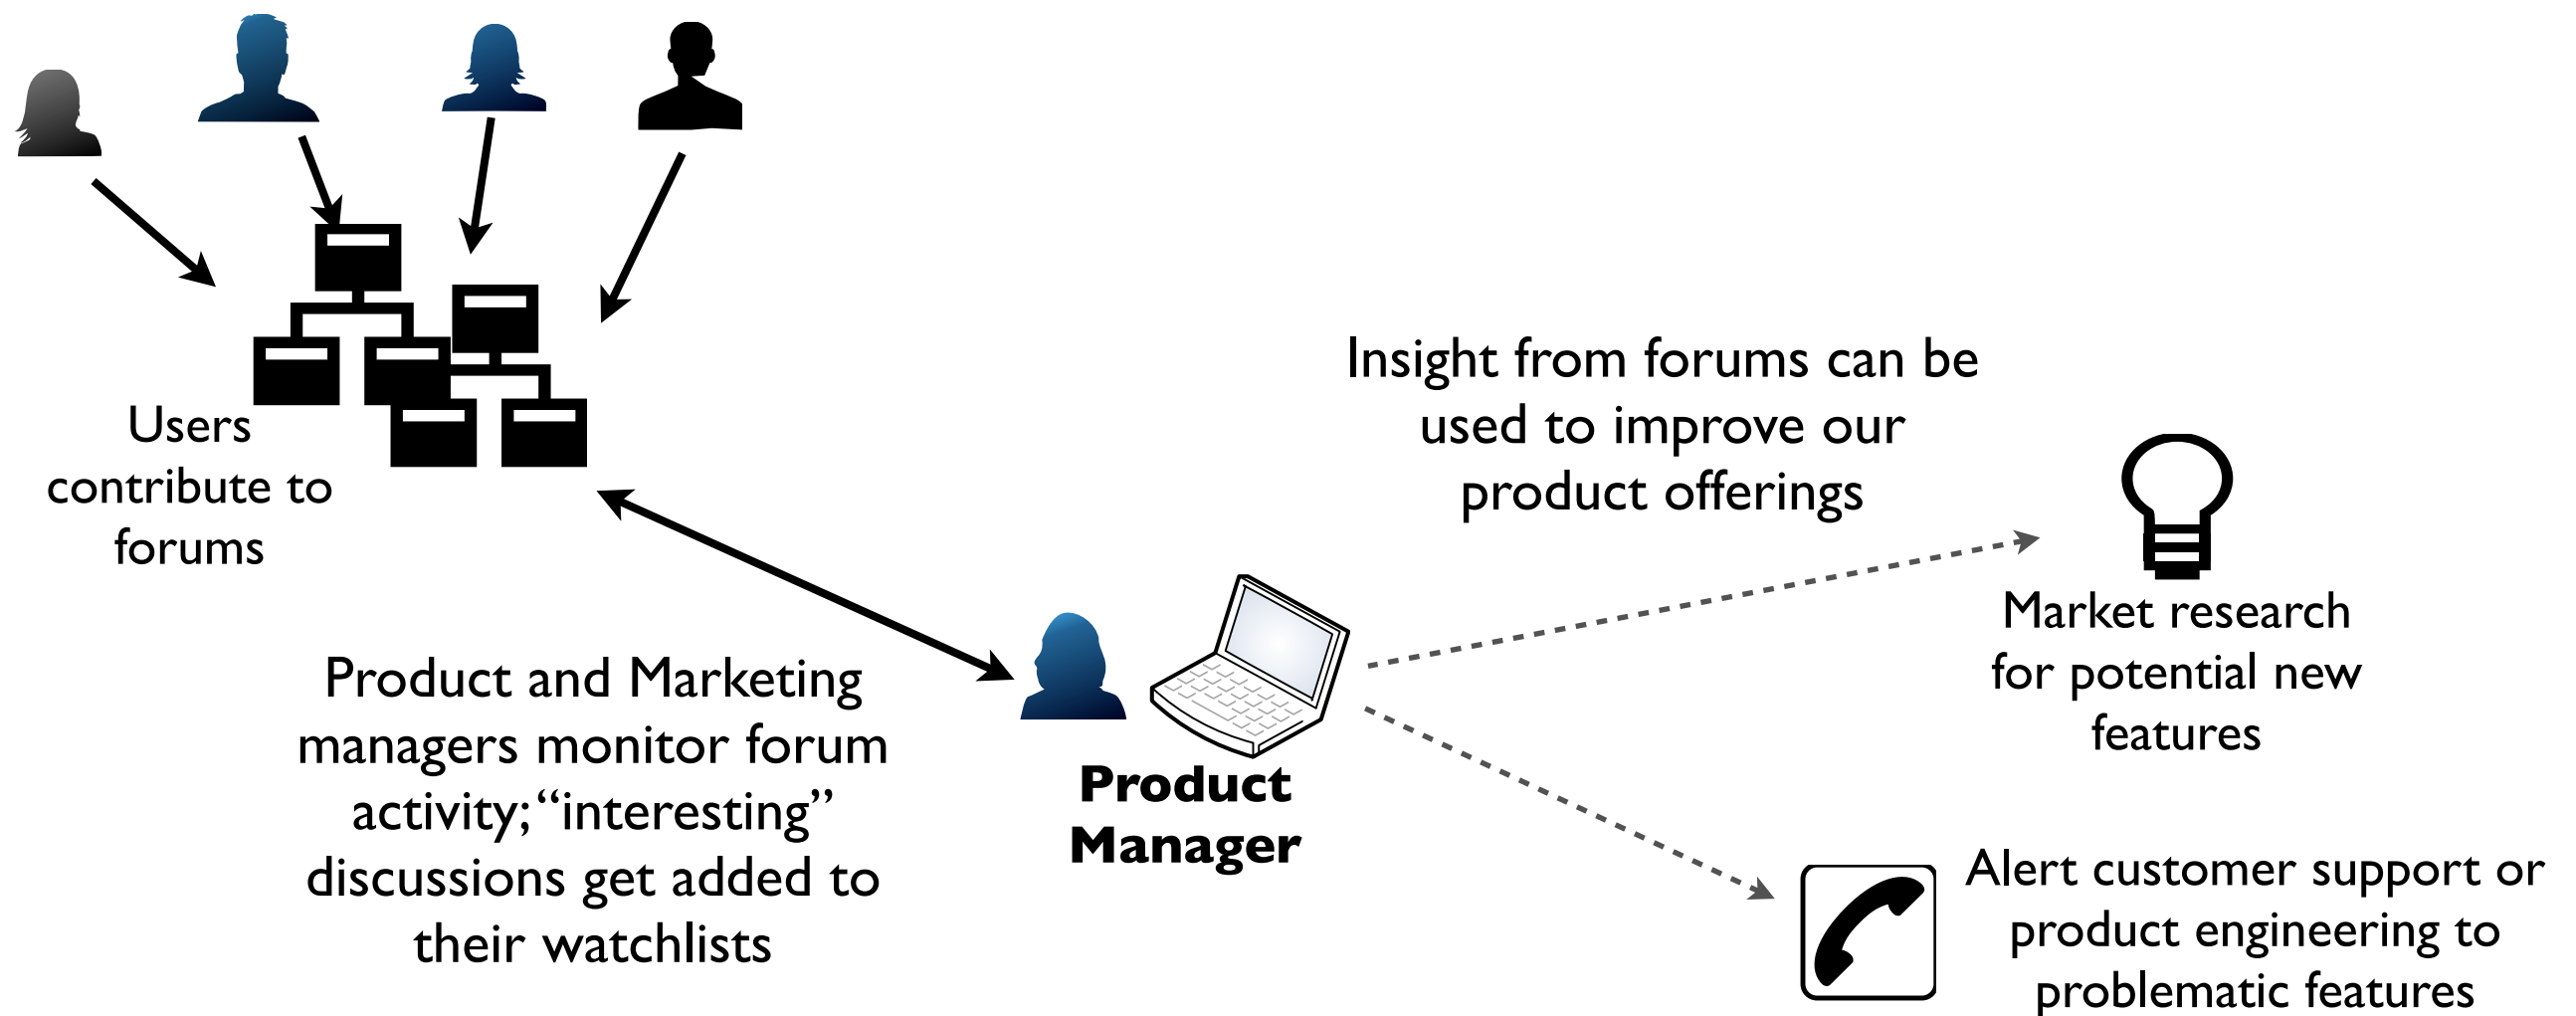
WVCorp Data Science Team
Notional Users Presentation

Our Goal:

Catch User Needs Early

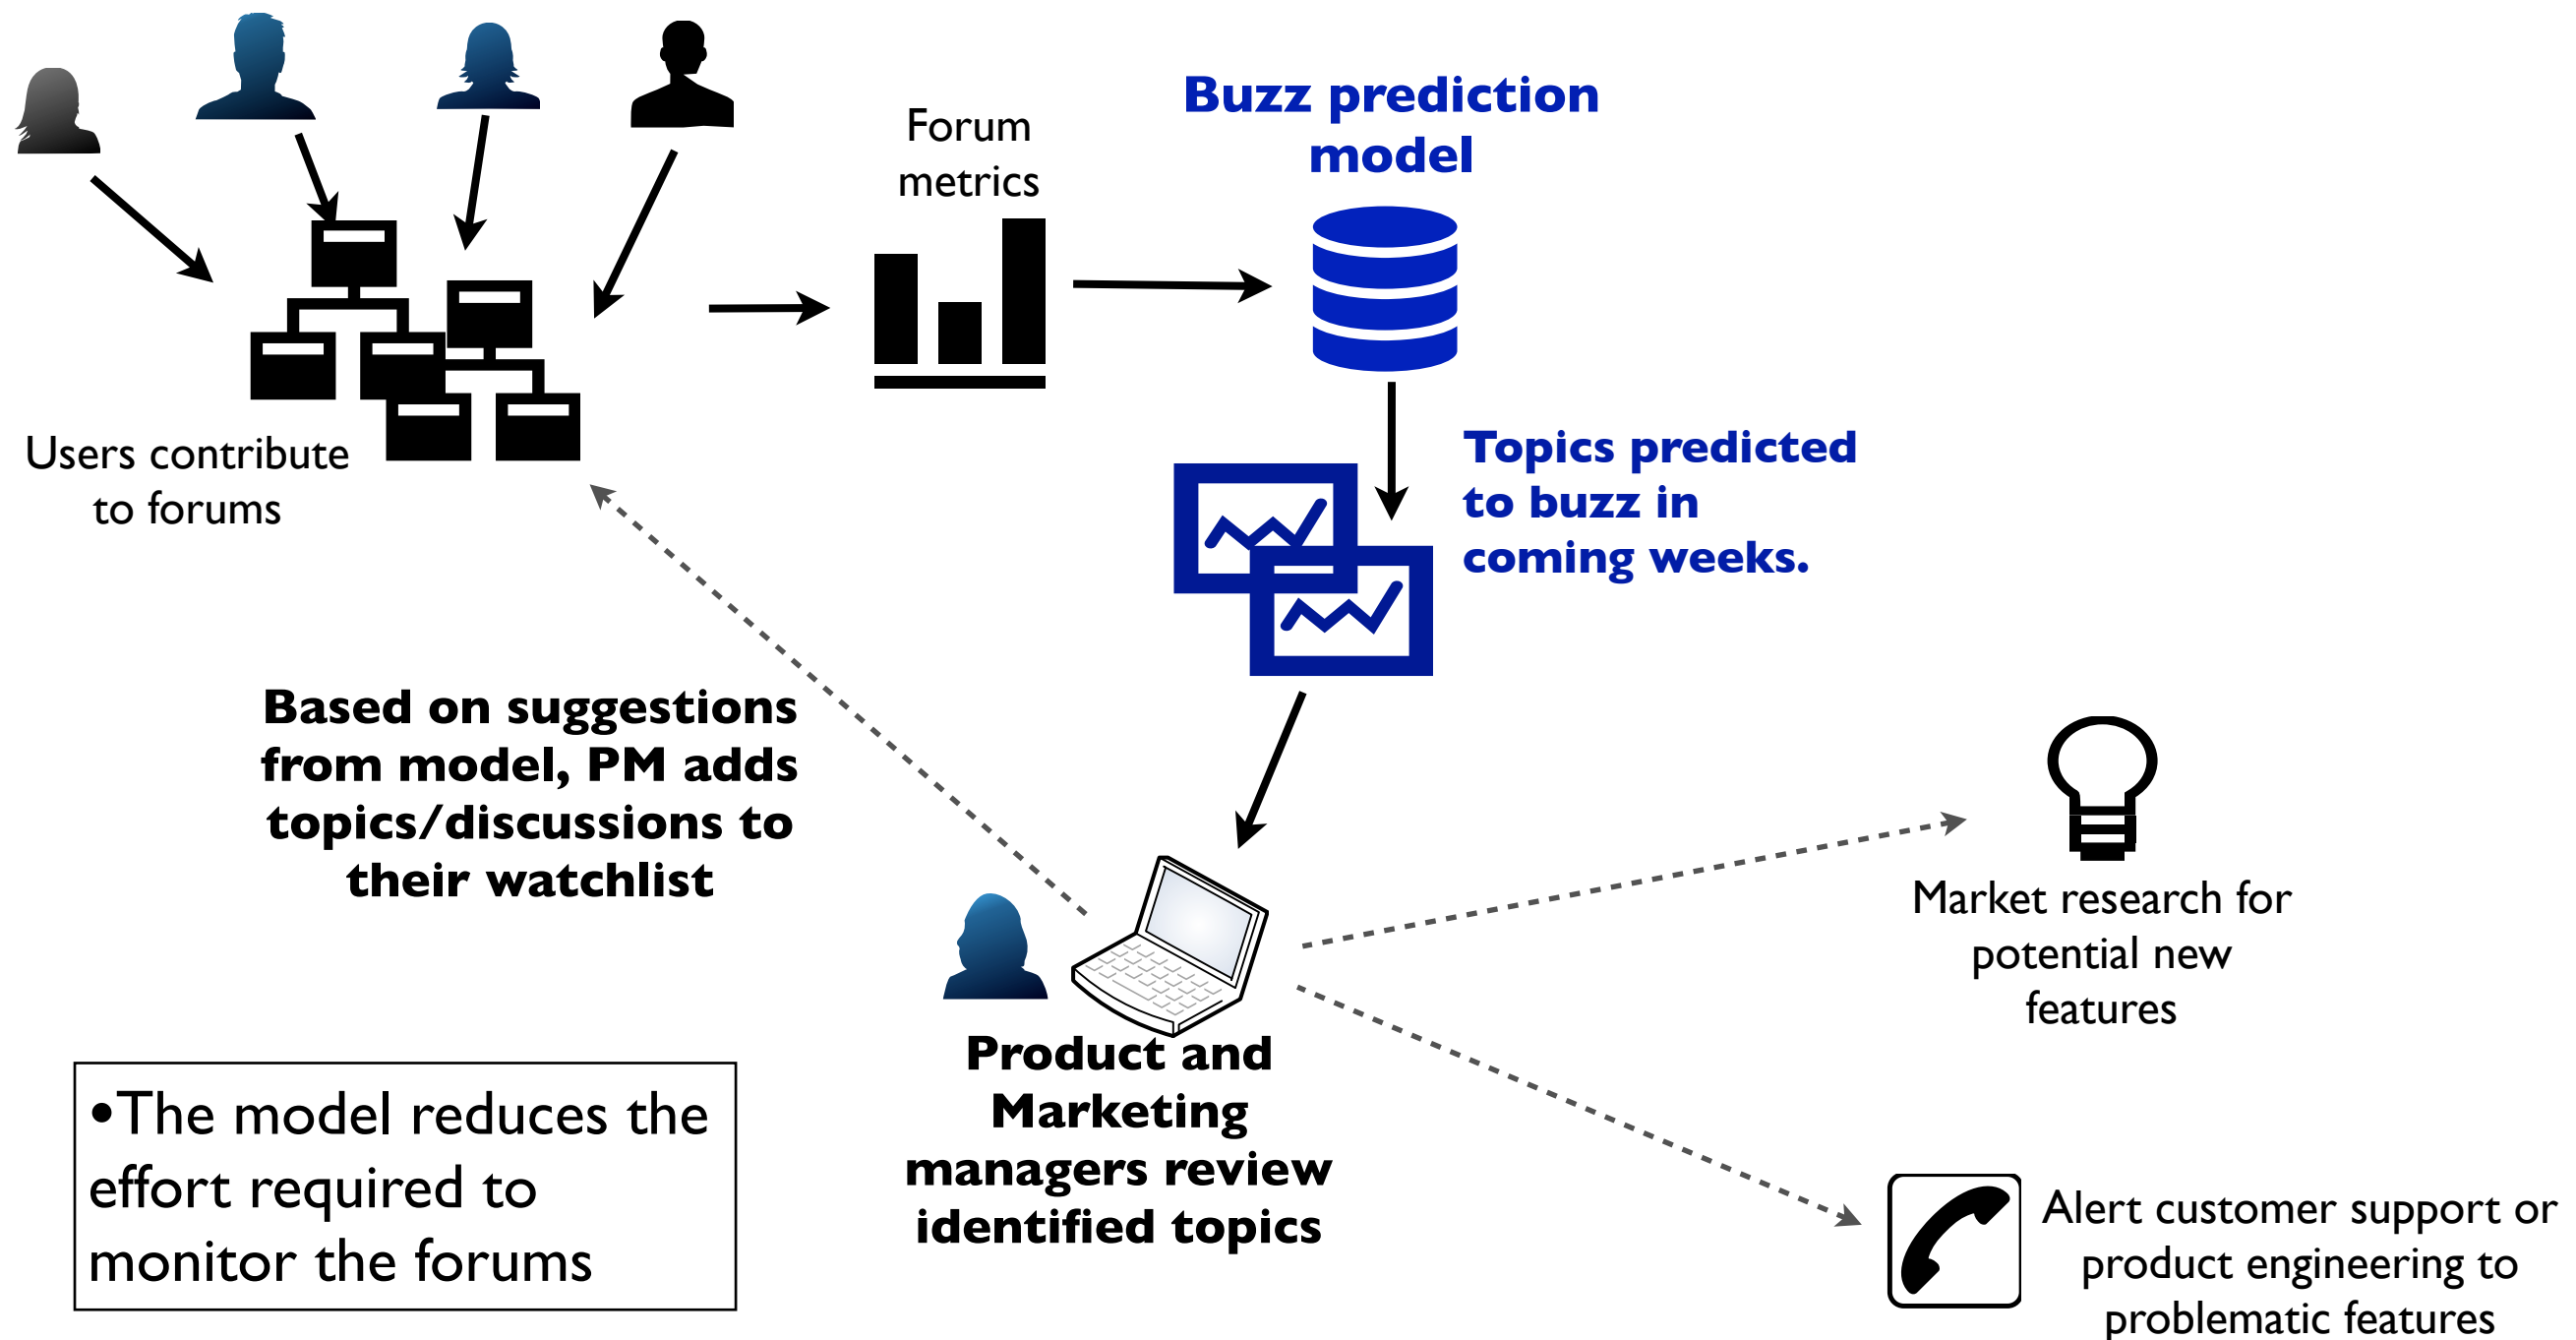
- Predict which topics on our product forums will have persistent buzz
- Features customers want
- Existing features users have trouble with
- Persistent buzz, not ephemeral or trendy issues
- Persistence = real, ongoing customer need

The Way it is Now



- Manually monitoring forums (even with watchlists) is time-consuming.
- Hundreds of topics, new discussions added every day

With The Buzz Prediction Model



... to how their workflow is changed (improved) by the model.

Find Information Faster

- Pilot Study: Reduce effort to monitor forums by a factor of 4
- Scan 184 topics -- not 791!
- Found 84% of about-to-buzz topics
- 75% of identified topics produced “valuable insight”

topics predicted to buzz that didn't

	Predicted No Buzz	Predicted Buzz	
No Buzz	579	35	614
Buzz	28	149	177
Total	607	184	791

about-to-buzz topics that were missed

topics the PMs can skip

topics the PMs have to review

Example:

Catching an Issue Early

- Topic: TimeWrangler → GCal Integration
 - # discussions up since GCal v. 7 release
 - GCal events not consistently showing up; mislabeled.
 - TimeWrangler tasks going to wrong GCalendar
 - **Hot on forums before hot in customer support logs**
 - Forum activity triggered the model two days after GCal update
 - Customer support didn't notice for a week

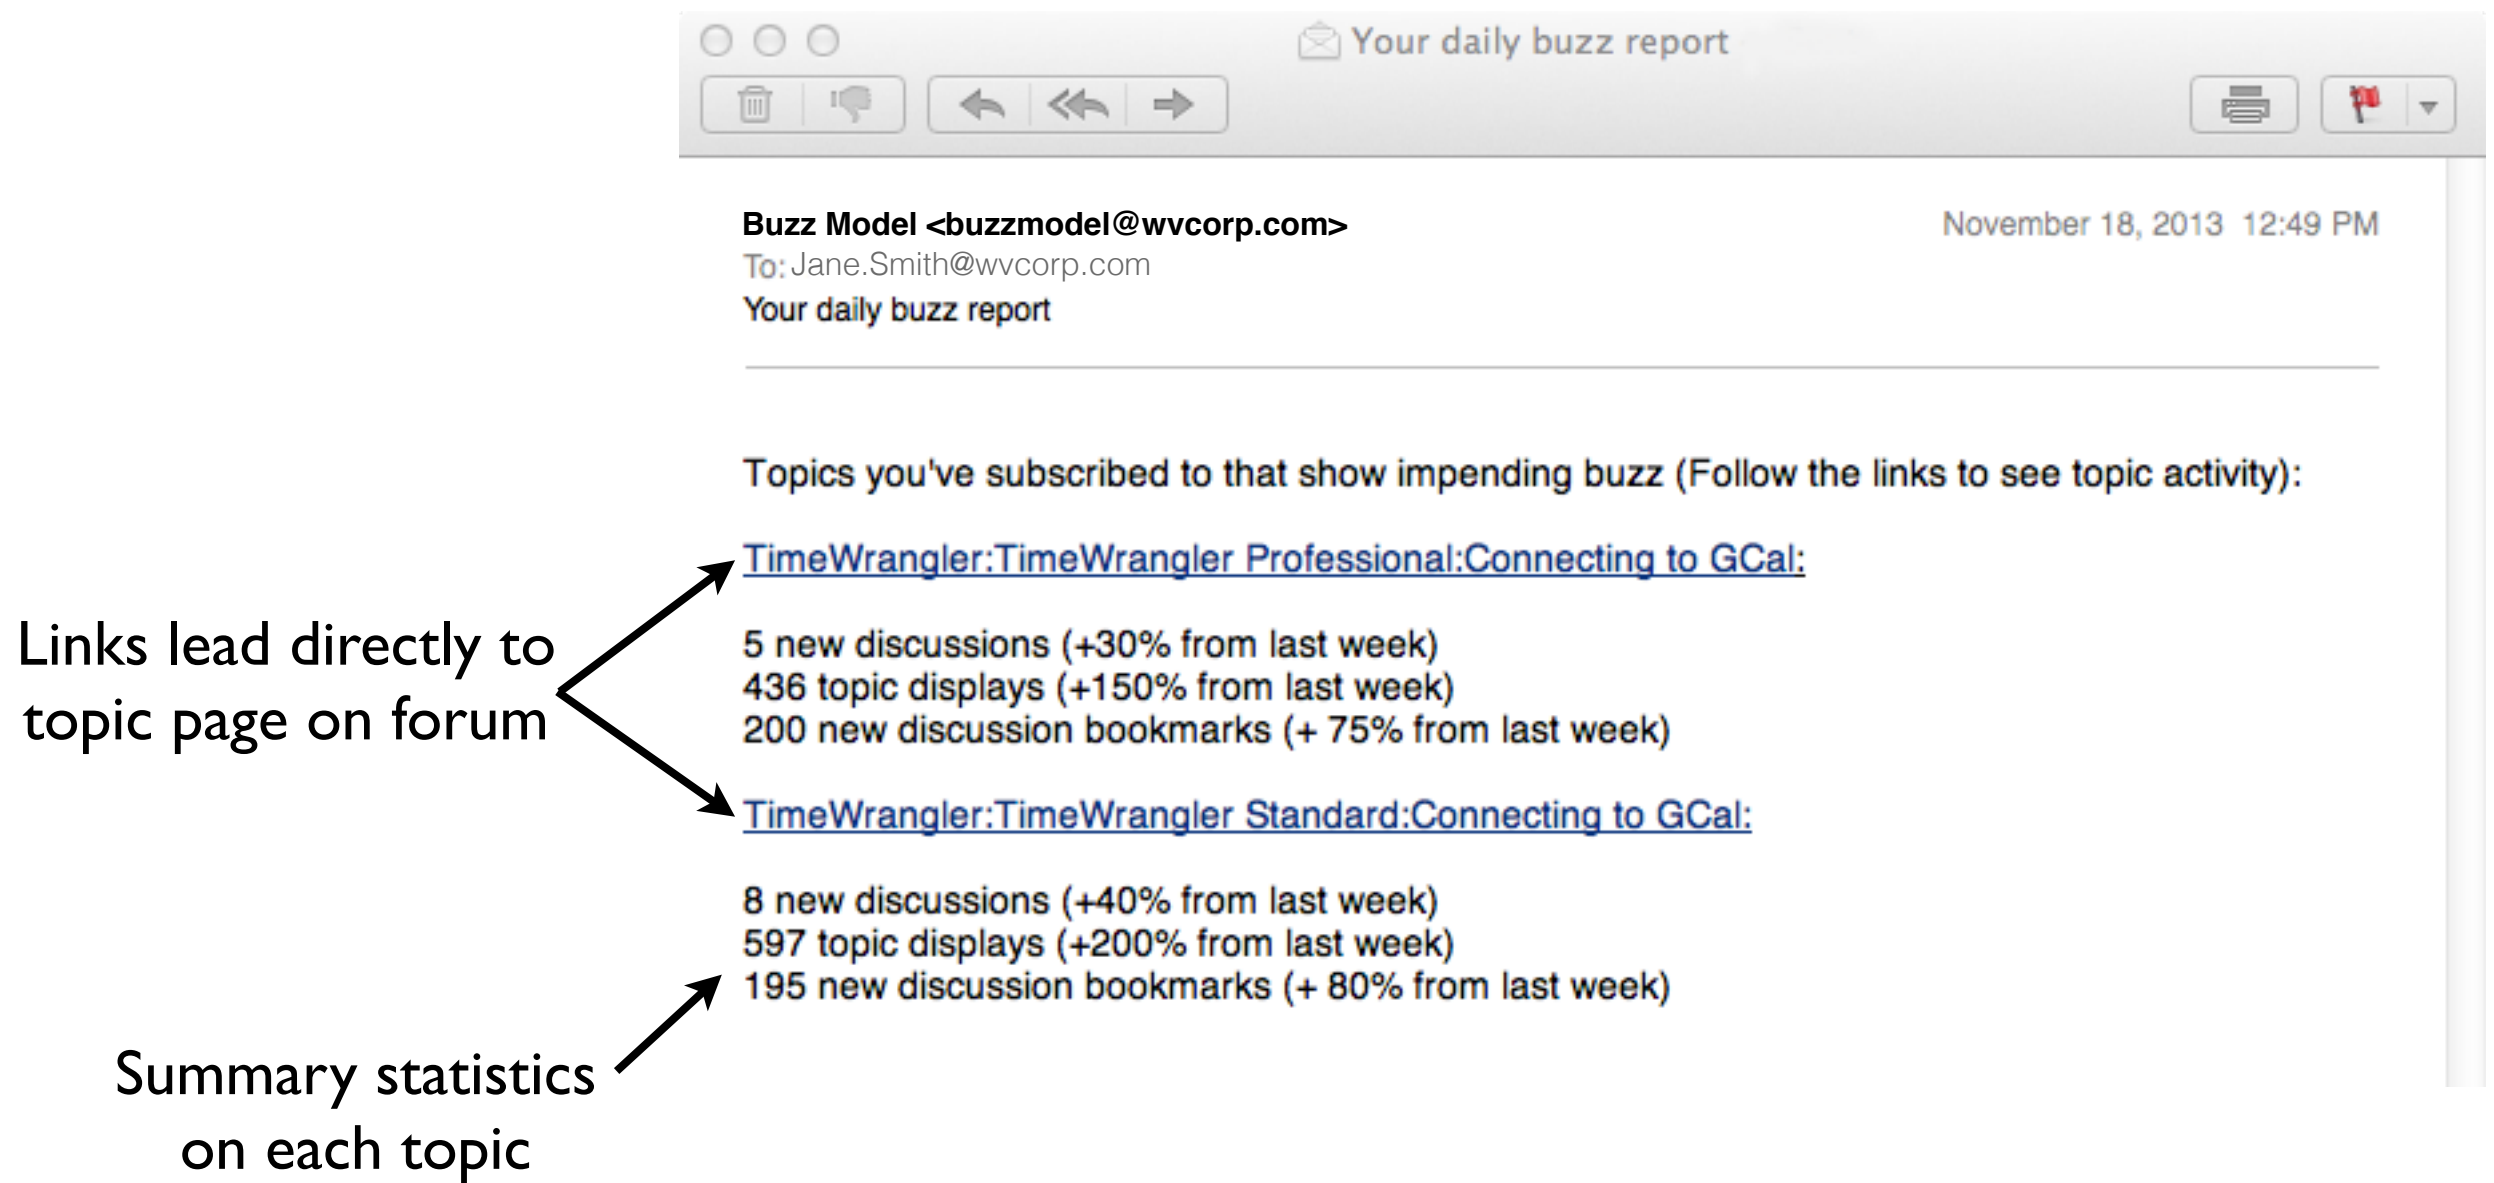
Metrics we Look At

- #Authors/topic
- #Discussions/topic
- #Displays of topic to forum users
- Average #contributors to a discussion in the topic
- Average discussion length in a topic
- How often a discussion in a topic is forwarded to social media

Using the Buzz Model

1. Go to <https://rd.wvcorp.com/buzzmodel> and register.
2. Subscribe to the product category or categories that you want to monitor.
3. Every day, the model will email you links to topics in your categories that are predicted to buzz (if there are any)
4. The links will lead you to the relevant topics on the forum
5. Explore!
6. Add topics/discussions of interest, to your watchlist, as usual.
 - We will monitor which topics you mark, to assess how effective our predictions are (how useful they are to you).

Step 3: Email Notifications



Your Feedback Will Help

- Better ways to get the information to you
 - Dashboard? Browser plugin? Is email fine?
- Additional metrics we might add to the model
- Advice on what is and isn't valuable. How can we better distinguish?
- Any other insight that comes from using the model

Thank You